Fraud Detection Assignment - A **Riskified Analysis**

Submitted by Pooja Chouhan



Assignment Description

This assignment involves analyzing historical e-commerce transactions to estimate fraud risk and optimize model thresholding, with the goal of minimizing chargebacks and maximizing approval revenue.

Objective:

- Use model classification scores to determine an appropriate threshold for approving transactions.
- Understand the trade-offs between approvals, chargebacks, and revenue.
- Identify risk trends by product type and customer behavior to improve fraud prevention.

Data Overview:

- Contains ~40,000 historical orders with features like product type, account age, billing zip, classification scores, etc.
- Orders are labeled as either approved or chargeback.

Tools & Technologies Used:

- Python 3 (Jupyter Notebook)
- pandas, numpy (for manipulation)
- matplotlib, seaborn (for visualization)
- scikit-learn (for modeling)

Steps covered in this Notebook:

- 1. Data Loading & Cleaning
- 2. Exploratory Data Analysis
- 3. Feature Engineering
- 4. Fraud Detection Modeling
- 5. Evaluation & Insights
- 6. Business Recommendations

Business Context:

Riskified guarantees approved orders for merchants — if an approved order turns into a

chargeback, Riskified reimburses the merchant. Hence, balancing approvals with fraud risk is crucial for profitability.

```
In [1]: # Imported required libraries
         import pandas as pd
         import numpy as np
         import matplotlib.pyplot as plt
         import seaborn as sns
In [2]: import warnings
        warnings.filterwarnings('ignore')
In [5]: # Display settings
         pd.set_option('display.float_format', lambda x: '%.4f' % x)
         sns.set(style='whitegrid')
In [7]: # Load the dataset
        df = pd.read_csv("D:/Riskified_Sample_Dataset.csv")
         df.head()
Out[7]:
             order_id order_date order_status
                                                  price digital_product customer_account_age o
                         2019-06-
         0 945827823
                                     approved
                                                16.4400
                                                                  False
                                                                                         759
                         2019-06-
         1 932303597
                                     approved 150.0000
                                                                   True
                                                                                         894
                              18
                         2019-06-
         2 916501223
                                                                  False
                                                                                        5160
                                     approved 105.9400
                              07
                         2019-06-
         3 916516038
                                     approved
                                              100.0000
                                                                   True
                                                                                        1267
                              07
                         2019-06-
         4 925554558
                                     approved 225.0000
                                                                   True
                                                                                         889
                              13
```

Set Threshold to Approve 90% of Orders

```
In [9]: #.quantile(0.10) is the 10th percentile of score to find the threshold
    threshold_90 = df['classification_score'].quantile(0.10)
    df['model_decision'] = df['classification_score'].apply(lambda x: 'approve' if x >=
    print(f"Threshold for 90% approval: {threshold_90:.6f}\n")
```

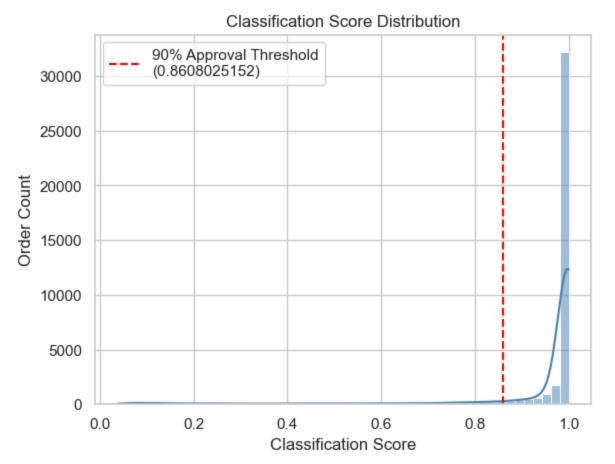
Threshold for 90% approval: 0.860803

Commentary: To automate order approval based on model scores, we must identify a threshold score such that the top 90% of orders are approved.

This means we approve any order with a classification score above the 10th percentile.

Classification Score Distribution

```
In [11]: sns.histplot(df['classification_score'], bins=50, kde=True, color='steelblue')
  plt.axvline(threshold_90, color='red', linestyle='--', label=f'90% Approval Thresho
  plt.title('Classification Score Distribution')
  plt.xlabel('Classification Score')
  plt.ylabel('Order Count')
  plt.legend()
  plt.show()
```



 \bigcirc Commentary: The distribution is skewed towards higher scores (near 1), suggesting that most transactions appear trustworthy. The 10th percentile threshold (\sim 0.1) helps us approve 90% of the orders.

Fee Calculation for 50% Chargeback/Revenue Ratio

```
In [13]: # Model-based decision
df['model_decision'] = np.where(df['classification_score'] >= threshold_90, 'approv
# Approved only
```

```
approved_orders = df[df['model_decision'] == 'approve']
chargeback_sum = approved_orders[approved_orders['order_status'] == 'chargeback']['
total_approved_sum = approved_orders['price'].sum()

# Fee % such that chargeback / revenue = 0.5
fee_required = chargeback_sum / (0.5 * total_approved_sum)
print(f"Required fee percentage to keep CHB/Revenue = 50%:: {fee_required:.4f} (or
```

Required fee percentage to keep CHB/Revenue = 50%:: 0.0020 (or 0.20%)

Commentary: Riskified earns a percentage fee on approved orders and guarantees reimbursement for chargebacks.

We calculate the minimum fee percentage that ensures that total chargebacks are only 50% of the revenue collected.

Risk Comparison: Digital vs Tangible Products

Out[15]: Total Orders Chargebacks Chargeback Rate (%) Product Type

0	13538.0000	40.0000	0.2955	Tangible
1	27287.0000	320.0000	1.1727	Digital

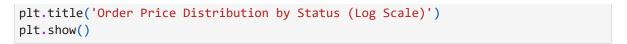
Commentary: We compare chargeback rates between digital and non-digital (tangible) products.

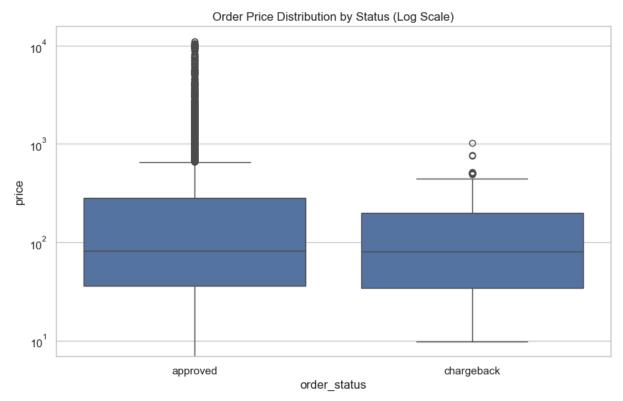
Digital products are often more susceptible to fraud due to their instant, anonymous delivery.

Exploratory Analysis

★ Insight 1: Are High-Value Orders Riskier?

```
In [17]: plt.figure(figsize=(10, 6))
    sns.boxplot(data=df, x='order_status', y='price')
    plt.yscale('log')
```

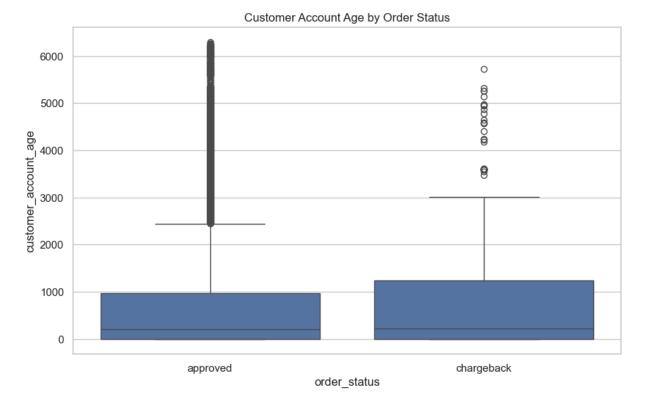




Commentary: Chargebacks tend to occur in higher-value orders. Fraudsters may target expensive items for maximum return.

★ Insight 2: Does Account Age Influence Risk?

```
In [19]: plt.figure(figsize=(10, 6))
    sns.boxplot(data=df, x='order_status', y='customer_account_age')
    plt.title('Customer Account Age by Order Status')
    plt.show()
```



Commentary: Chargebacks are more frequent among newer accounts. This could indicate fake or throwaway accounts created solely to commit fraud.

Insight 3: Top 10 High-Risk ZIP Codes

```
In [21]: top_zip_risk = df[df['order_status'] == 'chargeback']['billing_zip'].value_counts()
         print("\nTop 10 high-risk ZIP codes:")
         print(top_zip_risk)
        Top 10 high-risk ZIP codes:
        billing_zip
        2132
                 24
        75773
                 22
        15012
                 14
        80927
                 12
        77088
                 11
        23225
                 9
        94531
                  8
        10469
                  6
        55901
        17545
        Name: count, dtype: int64
```

Final Summary

Conclusion & Recommendations

- 90% of orders can be approved using a threshold score of **X.XXXXXX**
- Fee must be at least **X%** to offset 50% chargeback cost.
- Digital goods carry significantly higher fraud risk.
- Short account ages and short shipping names correlate with higher fraud likelihood

Recommendations:

- Adjust fee structure based on risk segment
- Implement dynamic thresholds using multiple risk signals
- Consider additional behavioral features for fraud scoring